

Programming Assignment 3: Instant Messenger

CSE 220: Systems Programming

Introduction

An instant messaging application allows users to communicate via text messages in real time. Users interact with the instant messaging application through a client. When a user composes a message and sends it from their client application, that message is sent to a server which then distributes it to all of the recipient client applications.

Communication between client and server is achieved by sending *packets*. The packets follow a specific format and carry a data payload in addition to other useful information about the message (such information is often called metadata).

In this assignment, you will implement the packet encoding and decoding for a basic instant messaging protocol. You will be using void pointers and pointer arithmetic, along with functions from `string.h`, to accomplish this. You will learn about pointer arithmetic, raw memory access, and data serialization. This assignment only requires you to encode and decode the outgoing and incoming packets, the rest of the instant messaging application is already implemented for you.

This assignment provides you with an instant messaging server and the remainder of the instant messaging application described in the previous paragraph. When the server is running, a client can connect and thereby communicate with other users also connected to that server. When the client process is started, the user's terminal window is split into two regions: a single line at the bottom where the user can type a command (the command region), and the remainder of the window at the top dedicated to the responses received from the server (the response region).

For example, if the user types the command

```
/stats
```

in the command region the following could be printed in the response region:

```
> CSE220_ServerBot - Statistics:
    Number of valid messages sent to server: 4
    Number of invalid messages sent to server: 0
    Number of refresh messages sent to server: 340
    Most active user: alphonse (4 messages)
```

Academic Integrity

As detailed in the course syllabus, academic integrity is important for the value and durability of your degree from the University at Buffalo. Here are a few reminders to help you maintain the integrity of this assignment.

- It is a violation of the course academic integrity policy to share this assignment document or details about this assignment with any student at UB not enrolled in CSE 220 during this academic semester, with any student at any other institution, or with anyone else without permission from your instructor. This includes homework- and note-sharing Internet sites such as Course Hero and Chegg.
- It is a violation of the course academic integrity policy to share *any code* from this assignment with any student at any time during or after this semester. You may discuss your code with course staff if necessary.
- It is a violation of the course academic integrity policy to seek assistance from other students or *any online resource not specifically approved by your instructor*. Forbidden resources include: Stack Overflow/Stack Exchange, GitHub/GitLab/BitBucket, ChatGPT or other AI assistants, students who have previously taken this course, etc.

- It is a violation of the course academic integrity policy to discuss implementation details with anyone except course staff.

Students found violating any of these policies, or any other academic integrity policy in the syllabus, will be sanctioned. Sanctions may include failure in the course or expulsion from the University. Make sure that you are familiar with the course academic integrity policies, as well as the policies of the department and University.

1 Getting Started

You should have received a GitHub Classroom invitation for this project. Follow it and check out the resulting repository.

As always, read over the entirety of this handout before starting the assignment. If you have not yet read Chapter 5 of *The C Programming Language* by Kernighan & Ritchie, you should also read that before starting. You will specifically find Sections 5.3, 5.4, and 5.5 to be very helpful.

You should read and understand all of the code in the `src` directory. You are not expected to read or understand any code in the `client` directory, although you may look at it if you wish.

Man page sections are noted as a number in square brackets after a keyword in this document. For example, `memset [3]` indicates that the manual page for the `memset()` function is found in section 3 of the Unix manual, and you can view it with the command `man 3 memset`.

2 Requirements

In this assignment, you must implement encoding and decoding of data packets. The packet formats used in this assignment are specified in Section 3.

You will implement four functions (note that each packet passed as a `void *`):

- `int pack(void *packed, char *input)`:
Parse input, a well-formed C string, to determine the packet type and fill `packed` with the encoded input.
- `int pack_refresh(void *packed, int message_id)`:
Fill `packed` with an encoded refresh packet according to the specification in Section 3.5.
- `int unpack(char *message, void *packed)`:
Decode `packed` based on the packet type and fill `message` with the decoded string.
- `int unpack_statistics(struct statistics *statistics, void *packed)`:
Decode `packed` according to the STATISTICS format specified in Section 3.4, filling `statistics` with the data.

Your program must parse user input and send the correct type of packet. User inputs that begin with special characters `"/` or `@` may be handled specially (see Section 3). No input that begins with `"/` that is not specified in this document will ever be tested, and you may use this to implement commands starting with slash for debugging (if it is useful to you).

For each of the above functions, the return value must be the integer value of the packet type as defined in *serialize.h*, or -1 for invalid inputs.

You should assume that any non-NULL pointer passed to a function is correctly allocated and indicates a region of memory of an adequate size. Other than that, you should not make any assumptions about the validity of inputs; you are responsible for validating them. This includes unreasonable or meaningless user input, as well as malformed data to be unpacked. Some examples of invalid values will be described in Section 3.

3 Packet Formats

The instant messenger application defines several packet types:

- **STATUS**
This packet type sets your status on the server. It should be sent when the user input starts with `"/me"` followed by a space. Any input starting with `"/me"` followed by any other character is invalid.
- **LABELED**
This packet type tags another user. It is sent when the user input starts with `"@"` followed by at least one and no more than `NAME_SIZE` non-space characters, followed by an ASCII space. `NAME_SIZE` is defined in `src/serialize.h`.
- **STATISTICS**
This packet type is used to request and receive a set of basic statistics from the server. It is sent when the user input starts with `"/stats"` followed by the end of the input. Any input beginning with `"/stats"` followed by any other character is invalid.
- **MESSAGE**
This packet type is just a normal text message sent by a user. It should be sent if the user input does not fall into any of the above categories.
- **REFRESH**
This packet is sent to the server by the client to request new messages to decode. This packet is different from the others in that it is sent once every second and *has no relation to user input*.

Note that the formats in this assignment are **very specific**, and must be implemented precisely. This includes details such as white space and the value of padding bytes. Be sure not to include any extra characters in quoted strings, insert extra newlines or other formatting, or deviate from this specification in any way!

The STATUS, LABELED, and MESSAGE packets have the following general form:

<code>int</code>	<code>char[NAME_SIZE]</code>	<code>size_t[]</code>	<code>size_t</code>	<code>char[]</code>
Packet Type	UBIT Name	Data Lengths	0	Data

The first part of the packet is common among the different packet types. It starts with an integer value indicating the packet type. This would be one of the packet types described above (REFRESH, MESSAGE, *etc.*), as defined in `serialize.h`. Following this is a character array of *exactly* `NAME_SIZE` bytes, (`NAME_SIZE` is also defined in `serialize.h`). This character array is valid if and only if it contains the sender's UBIT Name followed by sufficient ASCII NUL bytes to fill out the `NAME_SIZE` bytes of the array.

The remainder of the packet describes the actual data of the message. The "Data Lengths" field can be thought of as an array of type `size_t`, which varies in length based on the packet type. This list will always have a final value of 0, which terminates it. For each non-zero value in this list, there is an associated sequence of characters in the "Data" field. To give a simple example, if the Data Lengths field held the array {2, 5, 0}, and the Data field held the string "hihello", that would signify that there are two strings in the data area of the packet (of sizes two and five respectively), and that their contents are "hi" and "hello". Note that the data fields are not C strings, and do not include a terminating NUL byte! Note also that, due to ambiguity with the 0 terminating length value, *no valid packet can contain a field of length 0*. The sum of all data lengths in any packet must not exceed `MAX_MESSAGE_SIZE`, as defined in `src/serialize.h`.

The formats of the STATISTICS and REFRESH packets do not adhere to the general form given above, but can be found in the in-depth descriptions of each of the specific packet formats below.

3.1 MESSAGE

<code>int</code>	<code>char[NAME_SIZE]</code>	<code>size_t</code>	<code>size_t</code>	<code>char[]</code>
MESSAGE	UBIT Name	Message Length	0	Message

The message packet is the simplest data-carrying packet, holding only a single data field and its length. The data should be the entirety of the message input by the user *including any leading or trailing whitespace*. So if the user pagottes were to input "My name is Peter" the encoded message would look like this:

<code>int</code>	<code>char[NAME_SIZE]</code>	<code>size_t</code>	<code>size_t</code>	<code>char[]</code>
MESSAGE	"pagottes"	16	0	"My name is Peter"

Note that every character after the “s” in the UBIT Name field should be a NUL character, which you should remember has a integer representation of zero, and is idiomatically written as ‘\0’.

A message containing only ASCII space characters should not be encoded, and the pack function should instead return invalid.

When decoding this packet, the resulting string should be of the form “UBIT Name: Message”; the above packet would decode to “pagottes: My name is Peter”. Note that there is both a colon and a space between the UBIT Name and the Message.

3.2 STATUS

int	char[NAME_SIZE]	size_t	size_t	char[]
STATUS	UBIT Name	Status Length	0	Status

This packet type is very similar to that of the message type; however, the data it contains differs. As previously mentioned, the status packet type is sent when the user input begins with “/me” followed by one or more space characters. When you encode this input you should only send the status itself, not the “/me” or **any** of the spaces between it and the next non-space character. If the user input was “/me says hi”, the data encoded should be simply “says hi”; this would be the same for a dozen spaces of separation as a single space. Like the message packet type, a status of only spaces should be considered invalid.

A status should be decoded nearly the same as a message, but it should not have a colon, so it would be “UBIT Name Status”. If the user pagottes were to send a status resulting from the user input “/me says hi”, it would be decoded by a client as “pagottes says hi”.

3.3 LABELED

int	char[NAME_SIZE]	size_t	size_t	size_t	char[]	char[]
LABELED	UBIT Name	Message Length	Target Length	0	Message	Target

The labeled packet is the most complex type to pack. You will recall that this packet type is sent when the user input begins with an “@”, followed by one to no more than NAME_SIZE non-space characters, followed by a space. The characters between the “@” and the first space are the “Target”. Note that this does not include the initial character “@”. The characters starting with the first non-space character after the target is the “Message”. The same rules on space characters from the previous subsection apply. Any input starting with “@” that does not match the input format for a labeled message is an invalid input.

If the user pagottes input the text “@elb the Offspring is basically ska but good”¹, it would be encoded as:

1. The integer LABELED
2. The characters in “pagottes” padded out to NAME_SIZE bytes
3. A size_t containing 39 (the number of characters in “the Offspring is basically ska but good”)
4. A size_t containing 3 (the number of characters in “elb”)
5. A size_t containing 0
6. The 39 characters of the message
7. The 3 characters of the target

When decoding, the string is expected to be of the form “UBIT Name: @Target Message”. In the above example, it would be “pagottes: @elb the Offspring is basically ska but good”.

¹Peter’s musical discernment is known to be suspect. The Offspring is good, and they have some excellent ska tunes, but ska is king.

3.4 STATISTICS

int	char[NAME_SIZE]	char[NAME_SIZE]	int	long	long	int
STATISTICS	UBIT Name	Most Active	Most Active Count	Invalid Count	Refresh Count	Message Count

You should immediately notice that the statistics packet does not conform to the general format mentioned at the beginning of this section. Instead, it stores data corresponding to the members of the statistics struct defined in `serialize.h`:

```
struct statistics {
    char sender[NAME_SIZE+1]; /* Name of sender */
    int messages_count;        /* Number of packets sent to the server */

    char most_active[NAME_SIZE+1]; /* User who has sent the most messages */
    int most_active_count;          /* Number of messages sent by that user */

    long invalid_count;           /* Number of invalid packets sent to the server */
    long refresh_count;           /* Number of refresh packets sent to the server */
};
```

When encoding the statistics packet, none of this information is available to you, so you should only populate the packet type and your own UBIT Name. When decoding, you will decode this information into a statistics structure handed to the `unpack_statistics` function. The given client code will handle formatting the output printed to the screen.

3.5 REFRESH

int	char[NAME_SIZE]	int
REFRESH	UBIT Name	Last Message

The refresh packet is another packet which does not conform to the general case. It simply stores the packet type, your UBIT, and the ID of the last packet that your client received. This packet type is also special in the fact that you will only have to encode it, as it is never sent from the server to the client.

Because this packet type is used in the process of communication with the server, `pack_refresh()` is the first thing you should implement. If this function does not work correctly, you will not receive any messages.

4 Using the Client

When you start the chat client, you will be presented with a full-screen terminal application having a box at the bottom for user input. The chat client frequently calls `pack_refresh()` to send a REFRESH packet to the server to fetch incoming messages. Those incoming messages will be passed to `unpack()` to be displayed by your client. It will also accept user input and send that input to `pack()` when the user presses the Enter key.

To exit the chat client, press Control-C.

5 Testing

You are not given any direct tests for this assignment. Instead, you have the above specification and have been provided with a functional server to use for testing. The server will provide some feedback on the correctness of the packets that it receives. You should open two terminals, starting the server in one and the client in another. To start the server you will run `./chat_server server_sock`. To start the client, run `./chat server_sock`. The argument `server_sock` should be a path to *the same file*, which will be *created by the server*. Using a file that already exists will probably result in an error about "Address already in use"; if this happens, make sure there is not already a server running, and remove the file (e.g., `server_sock`). You should send at least one packet of every type from the client, and ensure that they are in turn displayed correctly on your client. You should avoid sending messages too quickly, however, as the server enforces a rate limit of 5 messages per second. If you

exceed the rate limit you will be locked out for 5 minutes, although when you are running the server locally you may just restart it.

You will find that this may be somewhat difficult to debug, as it can be hard to tell if the error is occurring in your encoding or decoding functions. For this reason, you will find `gdb` very useful for debugging. Of specific interest to you will be the `x` command. This command dumps memory at a pointer, as described by the `gdb` help message:

```
(gdb) help x
Examine memory: x/FMT ADDRESS.
ADDRESS is an expression for the memory address to examine.
FMT is a repeat count followed by a format letter and a size letter.
Format letters are o(octal), x(hex), d(decimal), u(unsigned decimal),
    t(binary), f(float), a(address), i(instruction), c(char), s(string)
    and z(hex, zero padded on the left).
Size letters are b(byte), h(halfword), w(word), g(giant, 8 bytes).
The specified number of objects of the specified size are printed
according to the format. If a negative number is specified, memory is
examined backward from the address.
```

Defaults for format and size letters are those previously used.
Default count is 1. Default address is following last thing printed with this command or "print".

You might use `gdb` to break at the end of your `pack` function and dump twenty bytes of memory in hexadecimal using `x/20xb` packed, for example.

Because the chat client uses the entire terminal window, running `gdb` and the chat client in the same window is inconvenient. Fortunately, `gdb` has the capability of connecting to another process on the system. You can do this by running `gdb filename PID`, where `filename` is the name of the binary to debug, and `PID` is the process ID of the running process. See the `gdb` help, [ps \[1\]](#), and the `gdb` resources provided for the `gdb` lab for more information.

Once you have a working implementation you will be able to connect to a shared, class-wide server on `emon.cse.buffalo.edu`. To connect to the shared IM server, build your project on Emon and do the following:

- `cd` into your git workspace for this project and run `make`.
- Invoke the client with `./chat /opt/chatserv/server/socket /opt/chatserv/client`. There is already a running server, so you should not start one.

Messages about failure to connect or bind, or Server Path: No such file or directory, probably mean that you made a typing error in the command. A message saying `bash: ./chat: No such file or directory` probably means that you are running the command from the wrong directory, forgot to run `make`, or your project failed to build.

Please note that the shared server will be monitored and logged, and you should absolutely not: behave inappropriately, use foul or offensive language, harass or disrespect anyone, or discuss any implementation details on it. This also applies any messages you send when you are locked out due to exceeding the rate limit.

6 Guidance

When considering how you are going to implement this assignment, we *highly* recommend reading over Chapter 5 of K&R. You will also want to read the `string [3]` man page, which contains a description of all functions defined in `string.h`.

You will almost certainly find it helpful to draw out a few encoded packets of each type, with the appropriate values and padding in the appropriate places, and think about how you will implement code to produce that output. Consider what input the user would type to produce the packet.

Learn how to attach `gdb` to your client process *sooner rather than later*. You will certainly find it helpful.

We recommend that, unless you have a different plan for identifiable reasons, you start by implementing `pack_refresh()`, and then move on to a version of `pack()` that handles *all user input as type MESSAGE* with little or no error handling, followed by `unpack()` capable of unpacking only MESSAGE packets. This will allow you to build out each stage using the given server (which will print errors if it receives invalid packets) and your own implementation (if you send a message, you should receive and unpack it with the expected results!) for testing. Once you have this minimal functionality in place, you can move on to the more difficult message types, error handling, and other tasks.

Remember to take time to design, as well as to refactor your code! If you find that you are copy-and-pasting, or reimplementing, more than a line or two of code, it may be time for a helper function. This project lends itself very well to a small number of helper functions that are called in several places each.

7 Grading

This assignment is worth 5% of your course grade. Points for this project will be assigned as follows; 10% of the assignment score for the handout quiz, and the remainder from this breakdown:

Points	Description
2	Handout quiz (will not be reflected on Autograder scores)
8	Packets are correctly encoded
6	Packets are correctly decoded
2	Inputs are correctly validated in <code>pack.c</code>
2	Inputs are correctly validated in <code>unpack.c</code>